

# **Judging an LLM on judging LLM-generated Microservice Documentation: A Supplemental Study and Observations**

Dawn Mai, 25R50028

Supervised by Takashi Kobayashi, School of Computing

# Agenda

01

## Introduction

Background, Motivation

02

## Methods

Proposed framework,  
working with repositories

03

## Results

Observation notes

04

## Conclusion

Summary of the study

05

## References

Related works

06

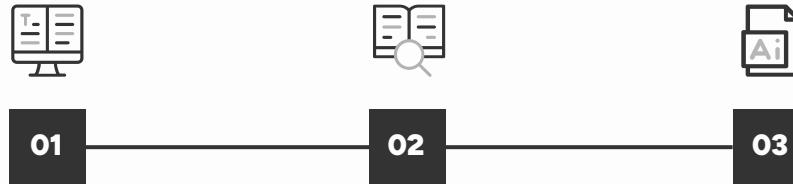
## Q&A

Questions?

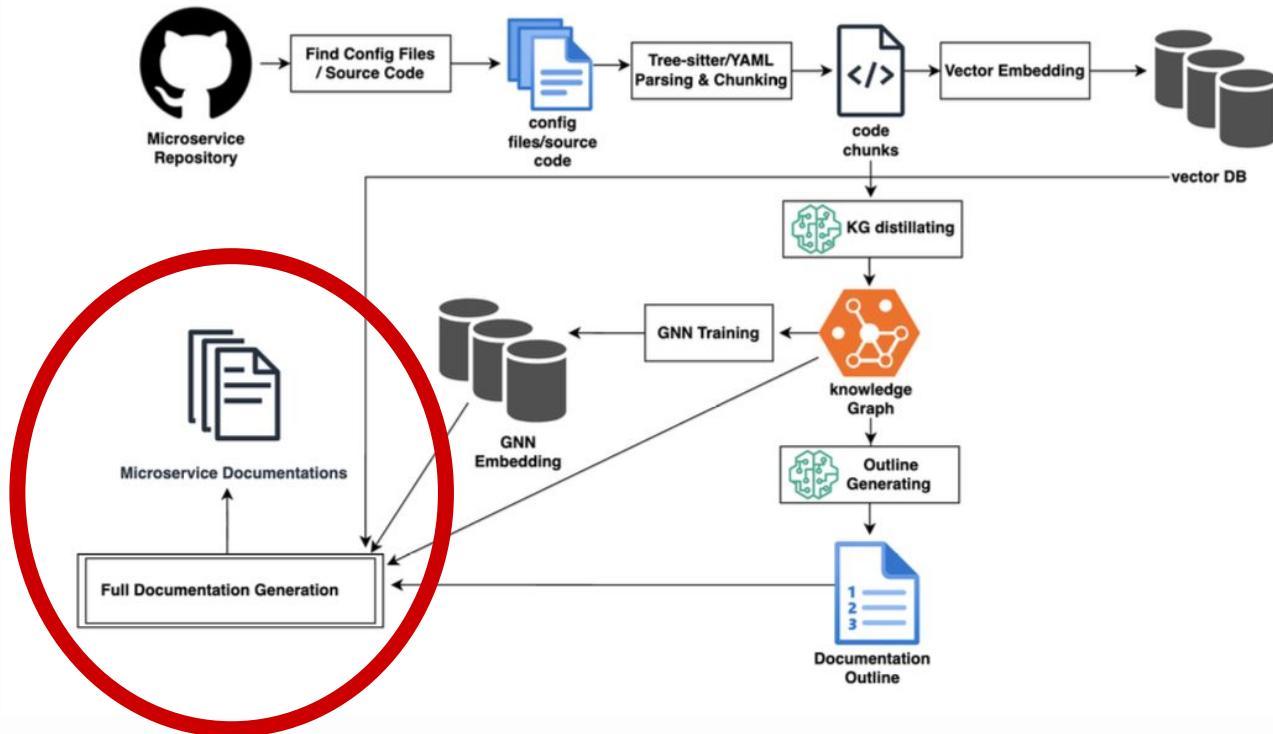
# Background

## Research goal

- Improve the pre-existing, original code from the provided thesis to correctly evaluate and rank the generated documentation for a new, partially-derived dataset from the original repositories.
- Enforce the ideal ranking:
  - KG/GNN → 3+
  - RAG → 1 - 3
  - Baseline → < 2



# Current Framework



Reproduced from QIAO Lin,  
Kobayashi Software  
Analytics Group [5]

# Issues (to fix in original code)

## # 1. Introduction

- The README file describes a food delivery app built using a microservice architecture.
- The technologies listed include Spring Boot, Spring Cloud, NestJS, Laravel, MySQL, PostgreSQL, MongoDB, Kafka, Zipkin, and Docker.
- The README does not explicitly mention any services named `gateway`, `ms\_order`, `ms\_payments`, or `service\_discovery`.
- The project is hosted on GitHub under the repository `JonasS6/food-delivery-app`.
- The app can be run using Docker, with instructions provided for using `docker-compose`.
- Contributions follow a "fork-and-pull" Git workflow.
- The project is distributed under the MIT License.

## # 2. High-Level Architecture

\*No info found (Graph Disconnected).\*

## # 3. Service Details

### ## `gateway`

\*No info found (Graph Disconnected).\*

\*No info found (Graph Disconnected).\*

\*No info found (Graph Disconnected).\*

## EVALUATION SUMMARY TABLE

filename	overall_score
documentation_rag_k5_a0.5.md	2.4
documentation_gnn_k5_a0.5.md	2.2
documentation_kg_ablation_kg_first_h2-1_k5.md	2
doc_0_baseline.md	1.4

# Experiment



## Fix the code!

Run it and manually debug + use Cursor to fix and explain why issues occurred and to adjust the LLM-as-a-judge prompts for more critical judgment



## Repository analysis

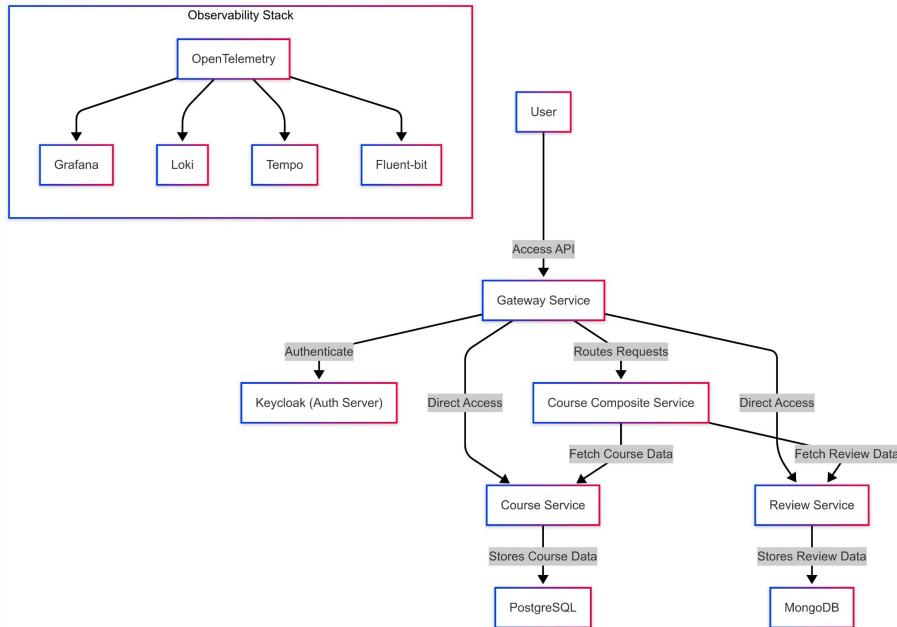
Three repositories from original thesis experiment + four new ones to further diversity the new dataset



## Document generation

Average scores of generated documentation was obtained by running the LLM prompt about 6 times; documentation of each attempt per repository was recorded on [Google Sheets](#)

# Experiment



Architecture diagram of  
Nasruddin/spring-boot-based-microservices

# Changes made:



**"No info found (Graph Disconnected)" error messages in the KG-ablation document:** this was fixed by forcing the framework to rely on semantic-based fallback and locate the graph instead of relying on an empty list.



**Low scores (scores of less than 3) given to KG- and GNN-based LLM-generated documents:** this was in part due to LLM hallucination of the presence/absence of services, which can result in documents saying a certain microservice was/wasn't present when the opposite was the case. This was fixed by tightening the "universal architect" prompt to penalize such contradictions, with a lighter penalty if the document in question was still useful. In addition, the "universal architect" prompts were edited to use a more strict, specific criteria, since we were looking for a certain ranking.



**RAG-based LLM-generated documents scoring higher than KG- and GNN-based documents:** the fixes in 3. were also applied here. No manual boosts were used to interfere with the data and neither the code generating these documents were altered; scoring only used the material of the documents provided and the repository they reported on.

# Improvement!



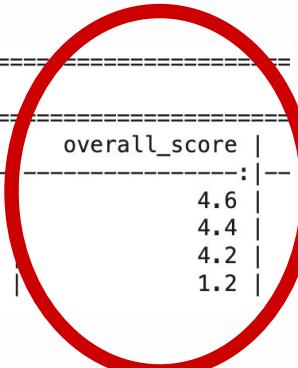
EVALUATION SUMMARY TABLE

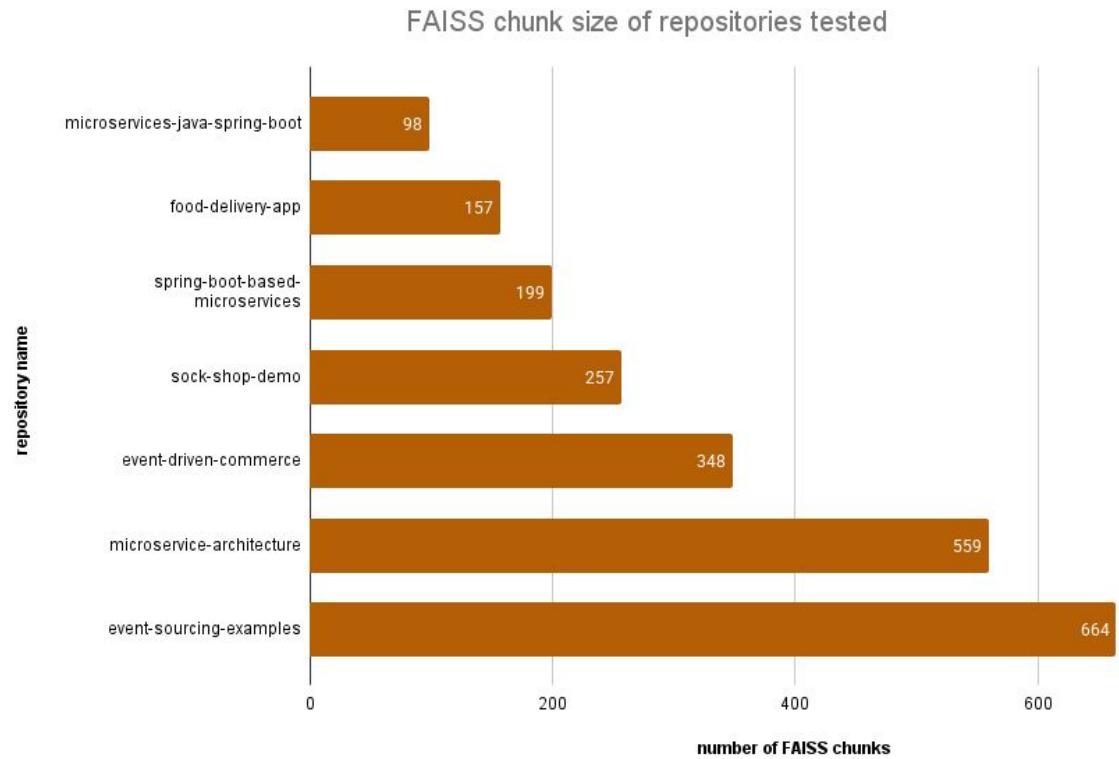
filename	overall_score
documentation_rag_k5_a0.5.md	2.4
documentation_gnn_k5_a0.5.md	2.2
documentation_kg_ablation_kg_first_h2-1_k5.md	2
doc_0_baseline.md	1.4



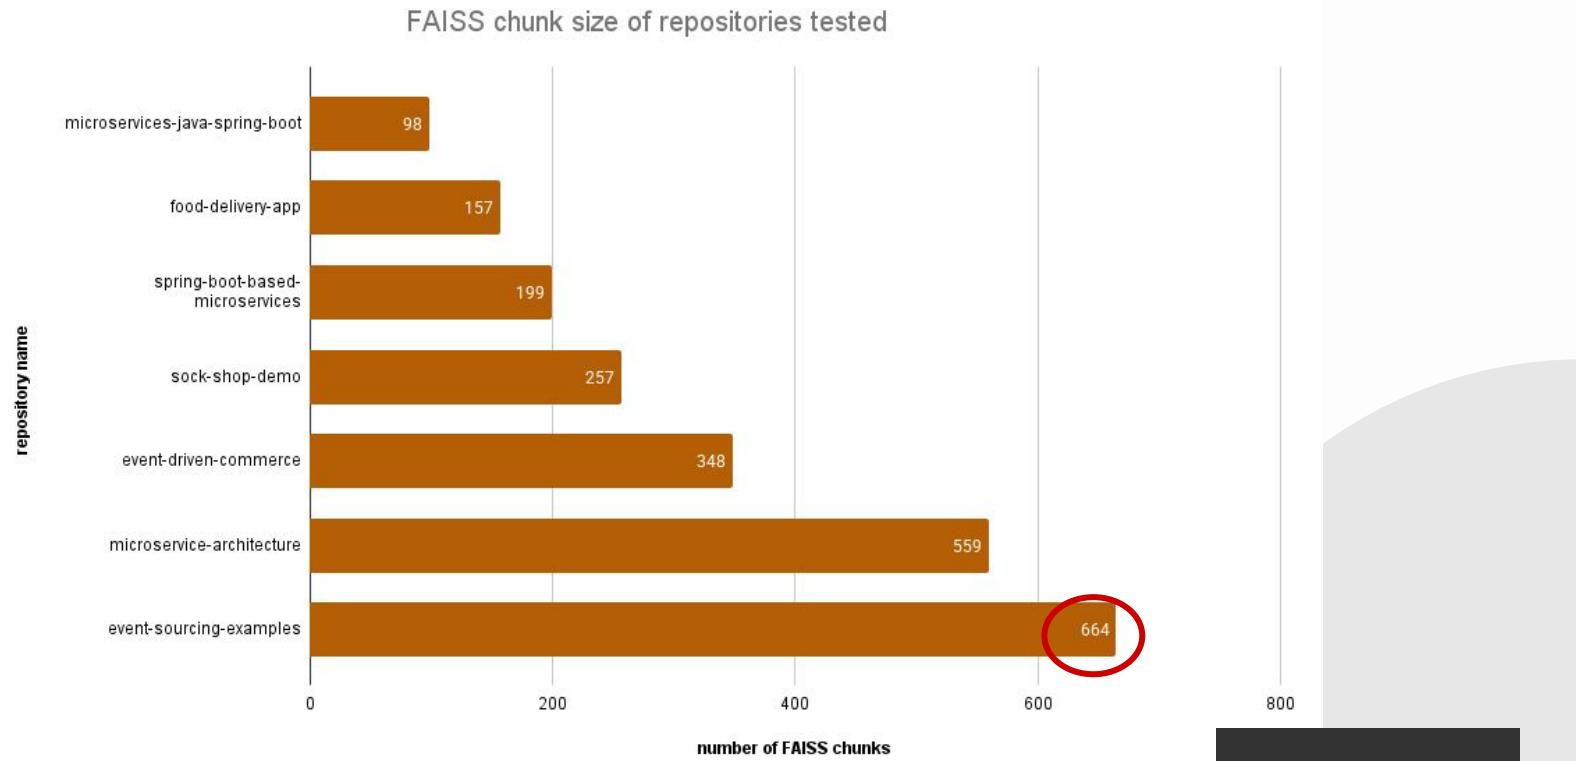
EVALUATION SUMMARY TABLE

filename	overall_score
documentation_gnn_k5_a0.5.md	4.6
documentation_kg_ablation_kg_first_h2-1_k5.md	4.4
documentation_rag_k5_a0.5.md	4.2
doc_0_baseline.md	1.2

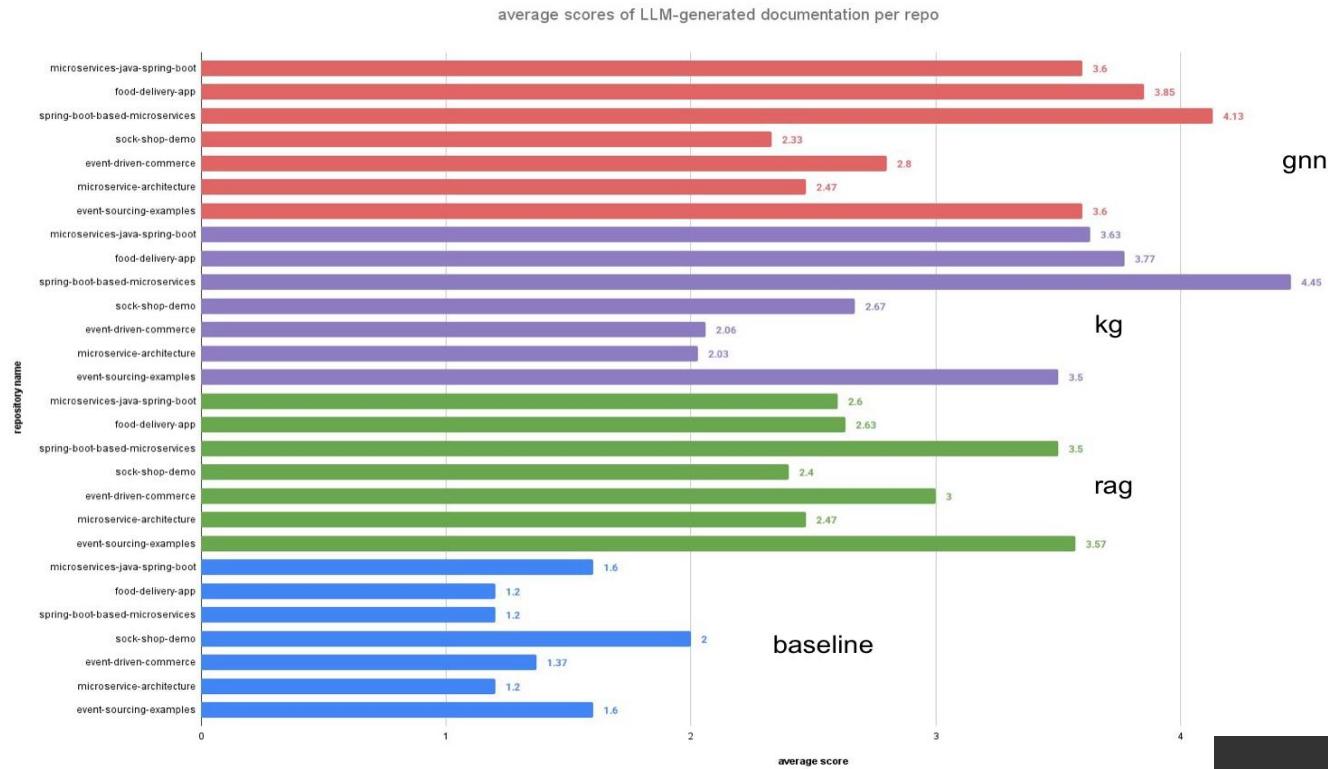




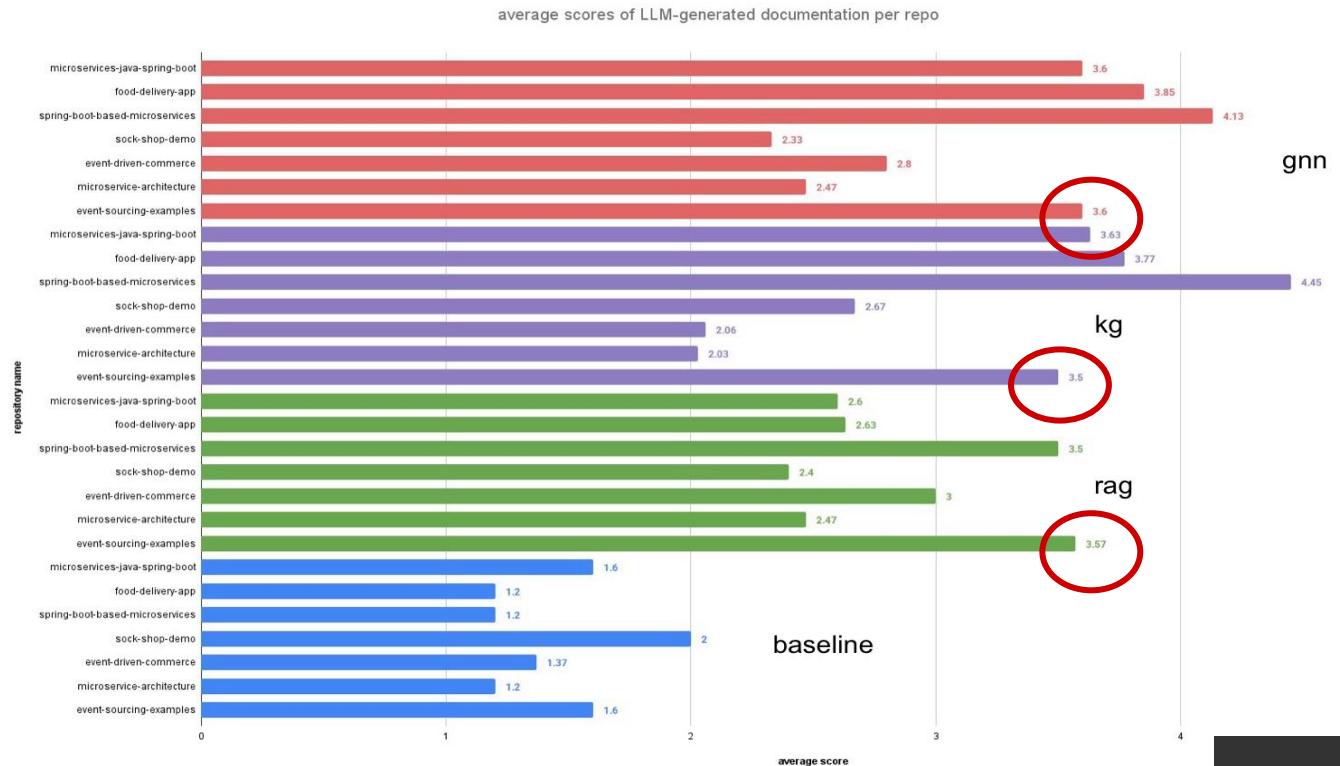
**Results**



**Results**



# Results



# Results

# Conclusion



## Observations

Did we improve the code? Yes!  
But we found that it works better on smaller microservices repositories.



## Limitations

Cost per OpenAI LLM call + one Neo4j instance = unable to test our framework on massive industrial scale repositories, LLM-as-a-judge bias



## Results

More code chunks → lower, less accurate scores to documentation; high scores for RAG documentation



## Future Work

Validate proposed method on more massive larger scale repositories + improve how ground-truth microservice architecture topology is extracted from a target repository [5].

# Works Cited

1. N. Dragoni et al., "Microservices: yesterday, today, and tomorrow," 2017.
2. J. Bogner, et al., "Assuring the evolvability of microservices: Insights into industry practices and challenges," Sept. 2019.
3. J. Bogner et al., "Microservices in industry: Insights into technologies, characteristics, and software quality," in 2019 IEEE International Conference on Software Architecture Companion (ICSA-C), pp. 187-195, 2019.
4. P. Lewis et al., "Retrieval-augmented generation for knowledge-intensive nlp tasks," 2021.
5. L. Qiao, "Structure-Aware Enhanced LLMs via Knowledge Graphs for Microservice Architecture Documentation", 2026. Master's Thesis, School of Computing, Institute of Science Tokyo.

# **Questions?**

# Thanks!

[dawnmai@cs.washington.edu](mailto:dawnmai@cs.washington.edu)

<https://github.com/dawnmai>



[www.linkedin.com/in/dawnmai](https://www.linkedin.com/in/dawnmai)

CREDITS: This presentation template was created by [Slidesgo](#), and includes icons, infographics & images by [Freepik](#)

Please keep this slide for attribution