

Grading an LLM on grading LLM-generated Microservice Documentation: A Supplemental Study and Observations

Dawn Mai, 25R50028

Supervised by Takashi Kobayashi, School of Computing

Agenda

01

Introduction

Background, Research
Goal

02

Methods

Proposed framework,
working with repositories

03

Results

Answering the research
goal

04

Conclusion

Summary of the study

05

References

Related Works

06

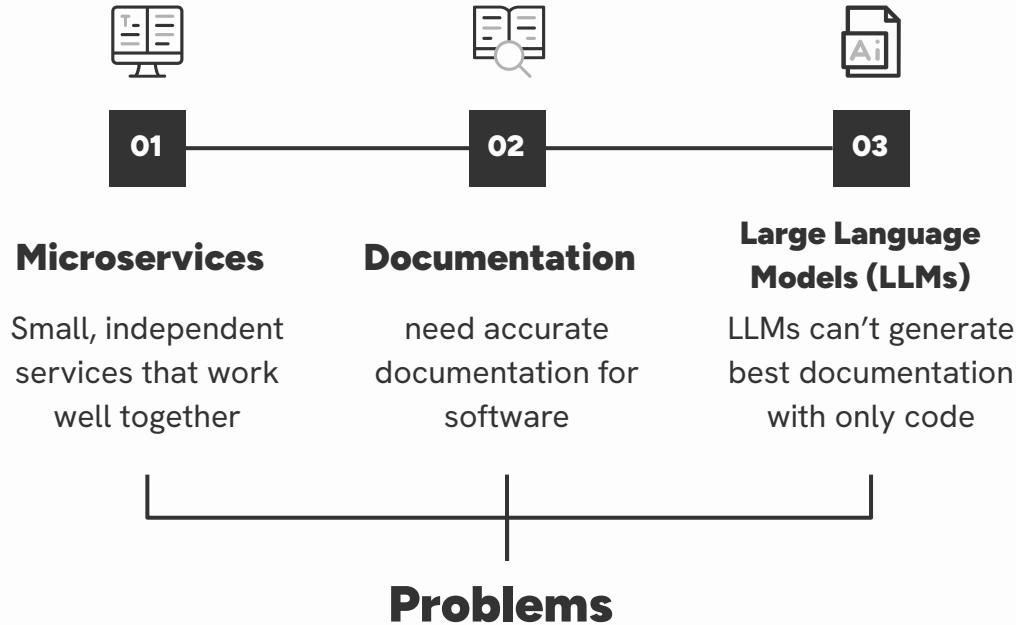
Q&A

Questions?

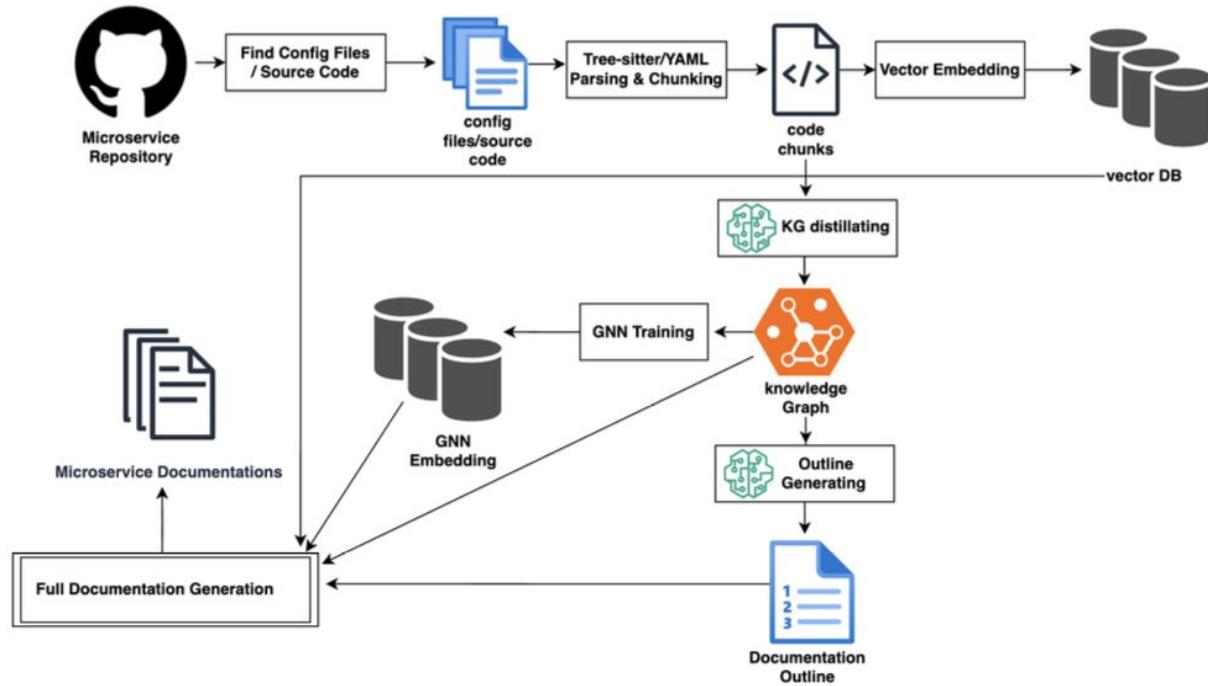
Background

Research goal

- provide structure to a given microservices architecture repository via constructing a Knowledge Graph (KG) and then a Graph Neural Network (GNN) based on said KG
- additional context to the LLM → enable the LLM to retrieve deeper context → generate high-quality documentation

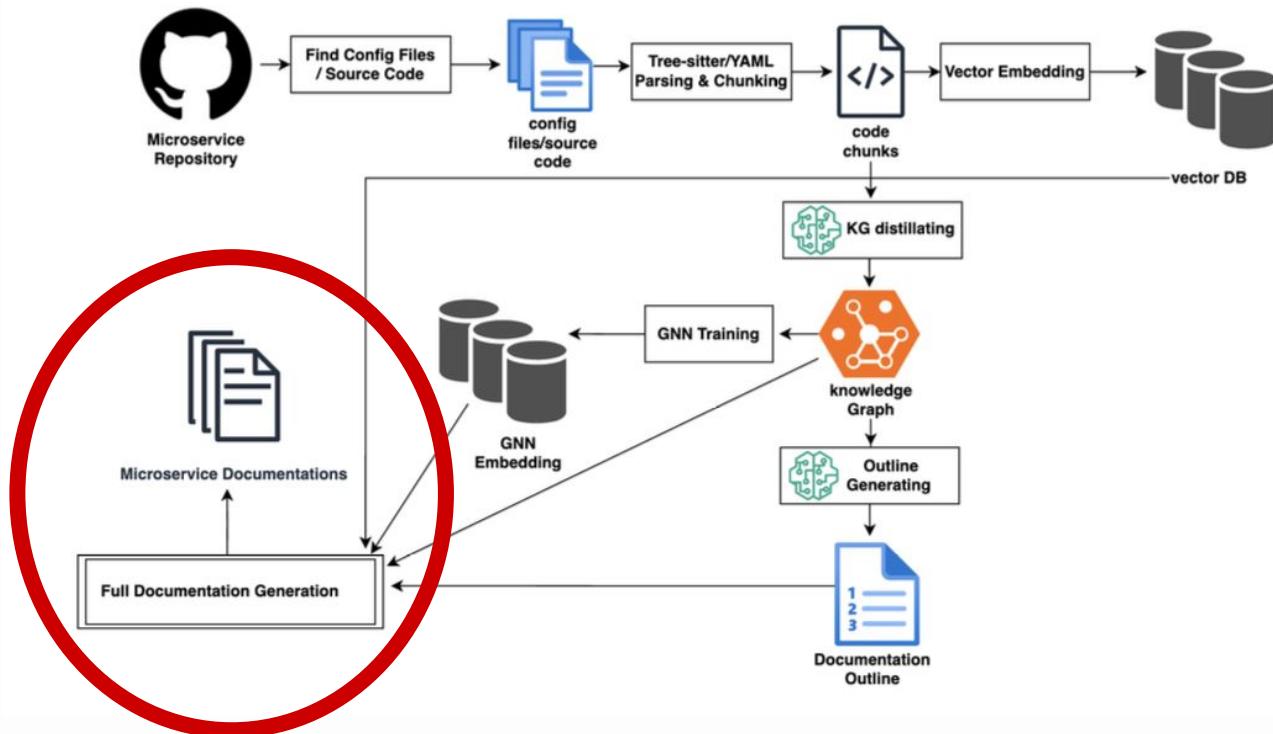


Proposed Framework



Reproduced from QIAO Lin,
Kobayashi Software
Analytics Group

Proposed Framework



Reproduced from QIAO Lin,
Kobayashi Software
Analytics Group

Experiment



Finding suitable repositories

Sourced off of GitHub and had to meet various requirements. Manual search as well as the use of AI tools (e.g. Cursor) were used to verify suitable repositories.



Repository analysis

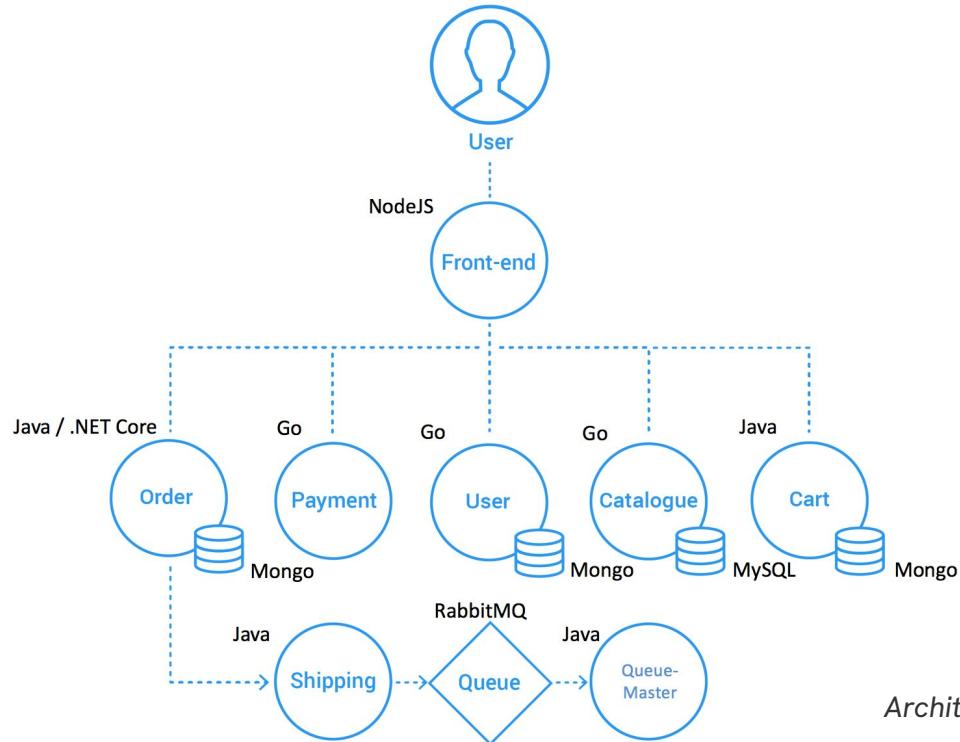
Code to build the KG for each repository was run once. Once the KG is built, code to build the resulting GNN and relative documentation is then executed.



Document generation

Average scores of generated documentation was obtained by running the LLM prompt about 6 times; documentation of each attempt per repository was recorded on Google Sheets.

Experiment



Architecture diagram of
[ocp-power-demos/sock-shop-demo](#)

Experiment

JUDGE_SYSTEM_PROMPT = """You are a Principal Software Architect acting as a Judge.
Compare the [GENERATED DOCUMENTATION] against the [GROUND TRUTH REPO].

****ABSOLUTE RULE:**** Use ONLY the provided ground-truth context (file structure, configs, service list).
Do NOT rely on prior knowledge of any repo or domain. If a claim is not supported by the ground-truth context, treat it as hallucination.

****EVIDENCE STANDARD:**** Prefer statements that can be directly verified from the provided context.
Penalize unsupported specificity more than cautious generality.
Penalize internal contradictions (e.g., claiming "no mention of X" while also describing X).

****SERVICE CONSISTENCY RULE:**** The provided service list is authoritative. Mentions of services not in that list should reduce ****faithfulness**** and ****correctness****. Missing most services should reduce ****completeness****.

****GENERALITY RULE:**** The rubric must apply to ANY repository. Do not assume microservices unless supported by ground-truth evidence.

**GLOBAL METRICS (1-5)**

||(1) COMPLETENESS**
* ****5:**** Covers ALL applicable sections with high detail from evidence.
* ****1:**** Misses major functional blocks that are clearly present.

||(2) CORRECTNESS**
* ****5:**** Details match ground-truth exactly.
* ****1:**** Hallucinates dependencies, ports, services, or versions.

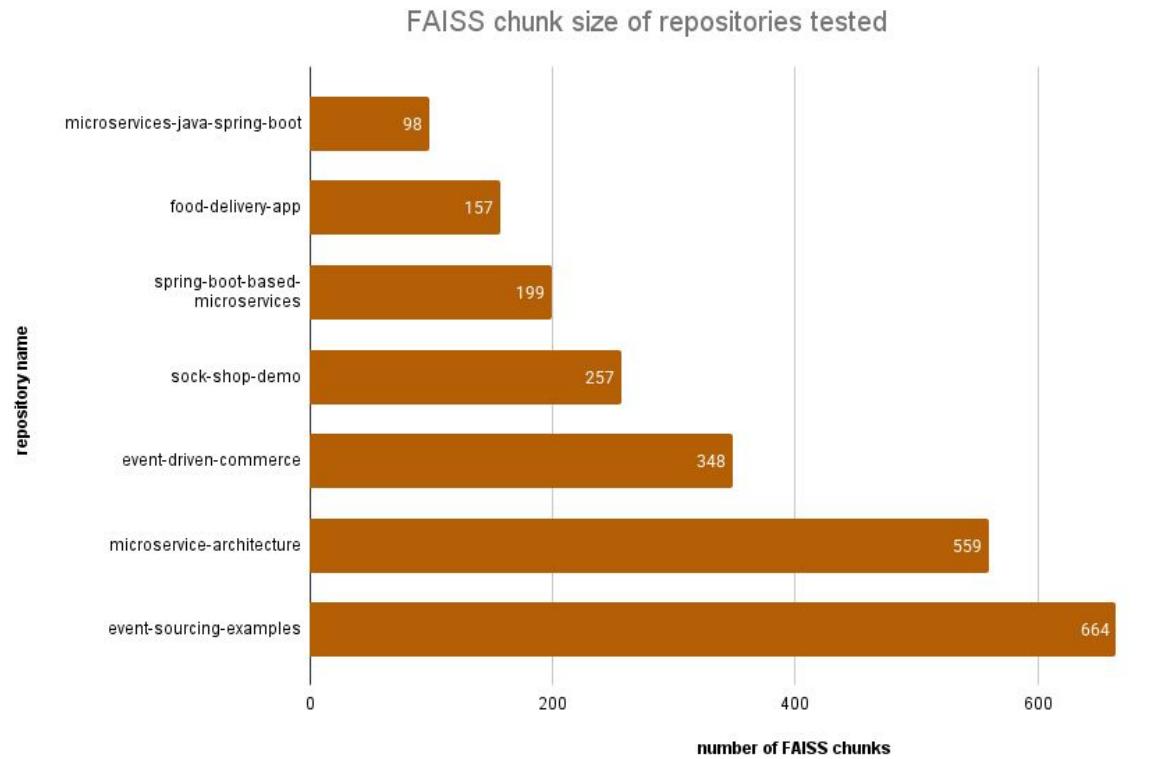
||(3) FAITHFULNESS**
* ****5:**** Claims are grounded in context.
* ****1:**** Invents functionality or entities not supported by context.

||(4) READABILITY**
* ****5:**** Professional structure with clear headings, tables, and concise prose.
* ****3:**** Understandable but inconsistently structured.
* ****1:**** Unstructured or hard to follow.

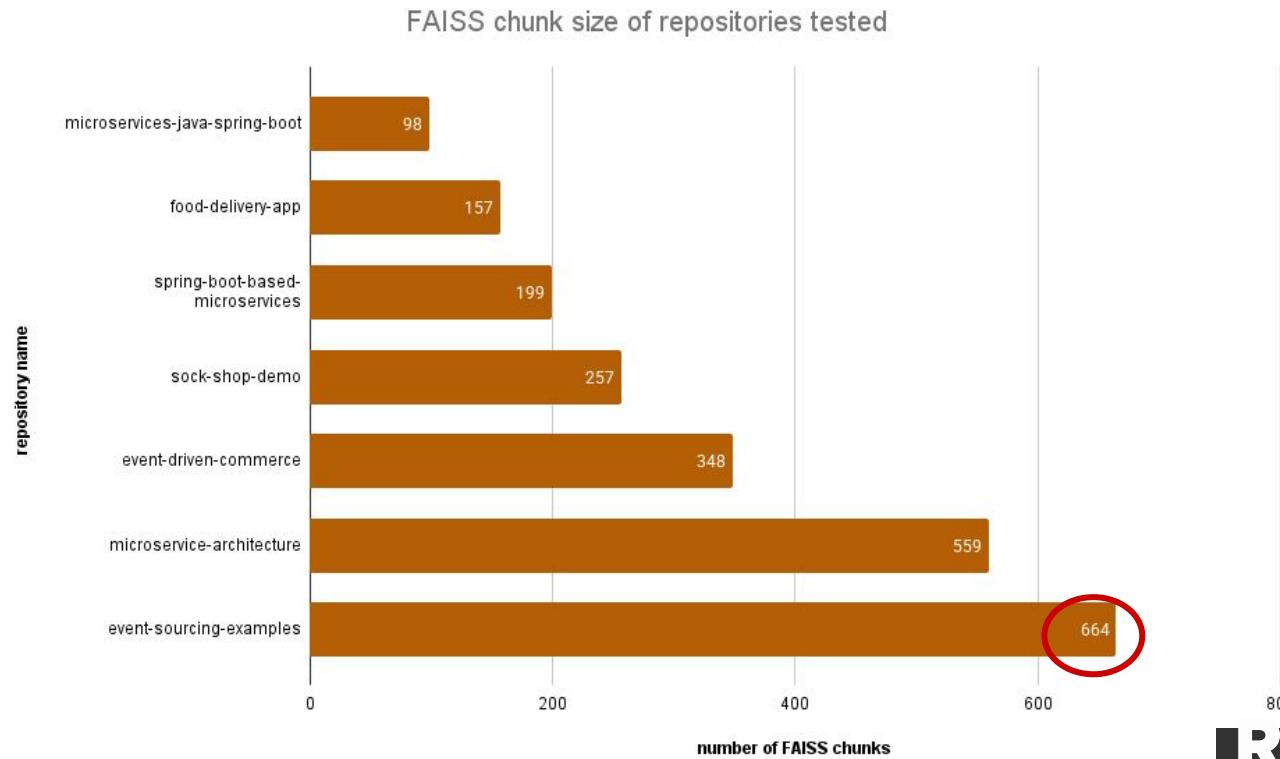
||(5) USEFULNESS**
* ****5:**** Actionable guidance derived from evidence (commands, configs, logs).
* ****3:**** Descriptive but not operational.
* ****1:**** Vague or marketing-like.



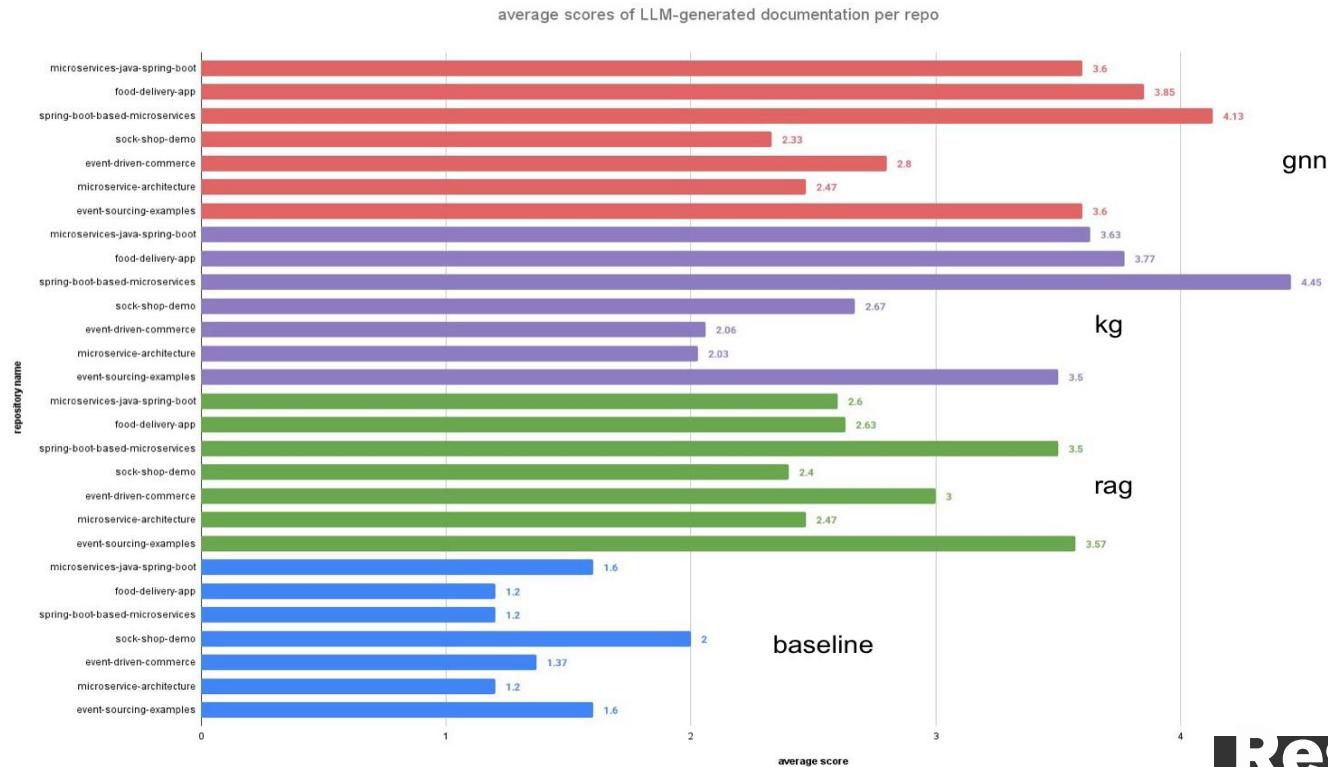
LLM-as-a-judge prompt for generated documents. Reproduced from QIAO Lin, Kobayashi Software Analytics Group



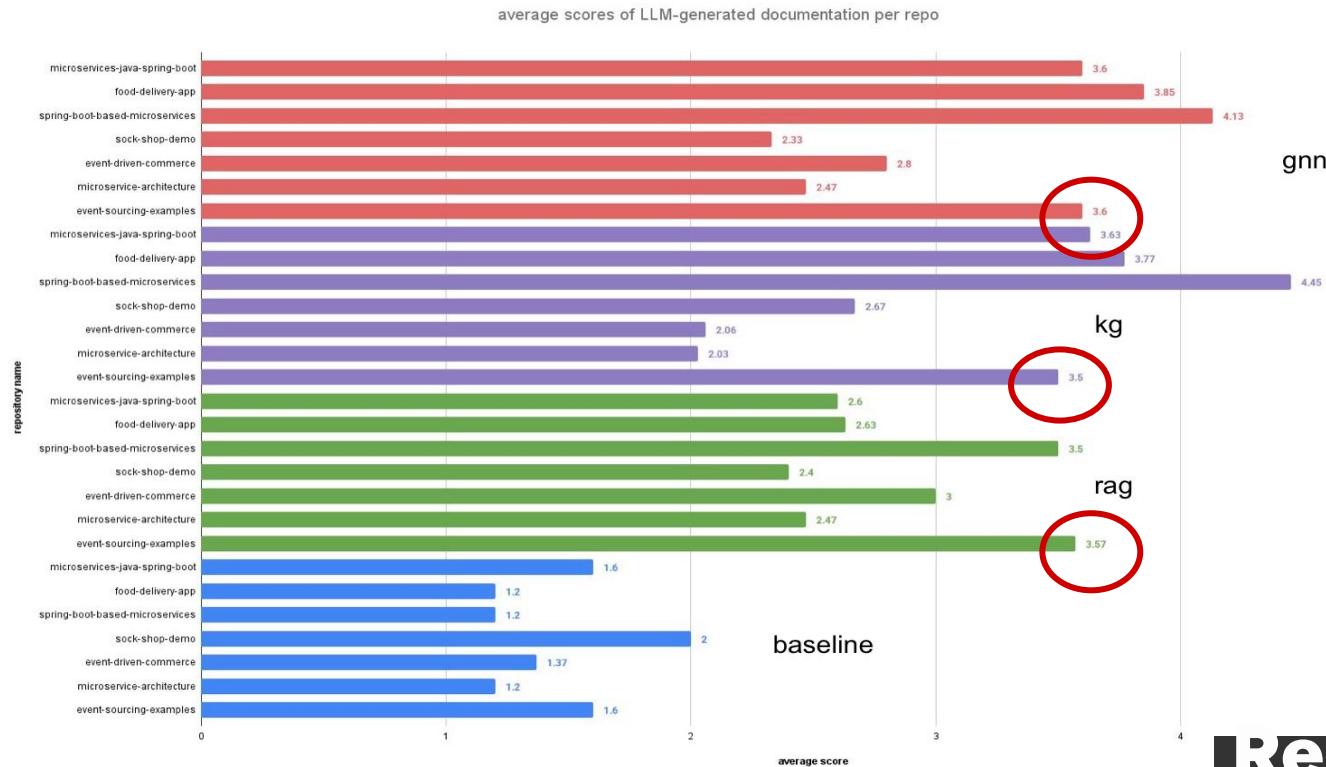
Result
s



Result
s



Result
s



Result s

Conclusion



Research Question

Can an LLM generate high-quality documentation given structure from code?



Limitations

Cost per OpenAI LLM call + one Neo4j instance = unable to test our framework on massive industrial scale repositories, LLM-as-a-judge bias



Results

More code chunks → lower, less accurate scores to documentation; high scores for RAG documentation



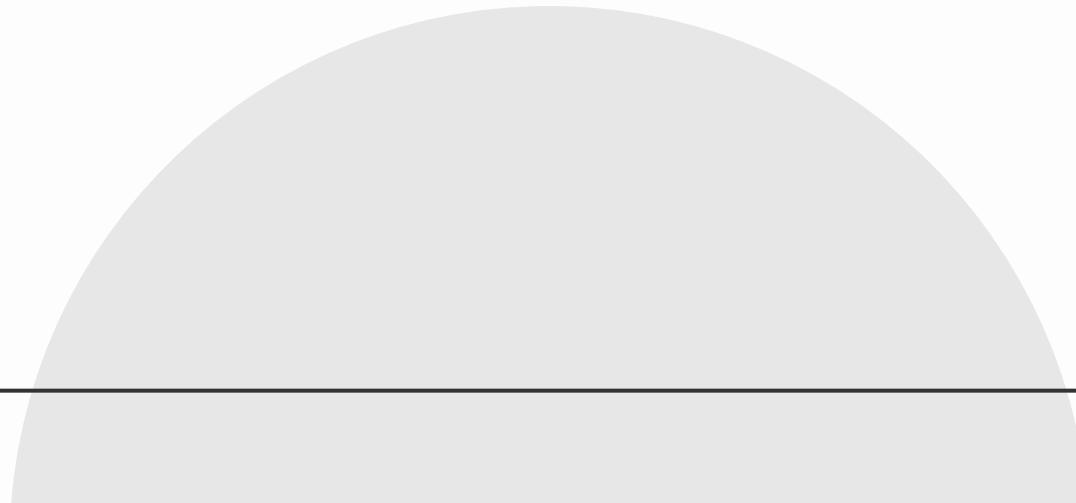
Future Work

Validate proposed method on more massive larger scale repositories + improve how ground-truth microservice architecture topology is extracted from a target repository

Works Cited

1. N. Dragoni, S. Giallorenzo, A. L. Lafuente, M. Mazzara, F. Montesi, R. Mustafin, and L. Safina, "Microservices: yesterday, today, and tomorrow," 2017.
2. J. Bogner, J. Fritzsch, S. Wagner, and A. Zimmermann, "Assuring the evolvability of microservices: Insights into industry practices and challenges," Sept. 2019.
3. J. Bogner, J. Fritzsch, S. Wagner, and A. Zimmermann, "Microservices in industry: Insights into technologies, characteristics, and software quality," in 2019 IEEE International Conference on Software Architecture Companion (ICSA-C), pp. 187-195, 2019.
4. P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. tau Yih, T. Rocktäschel, S. Riedel, and D. Kiela, "Retrieval-augmented generation for knowledge-intensive nlp tasks," 2021.

Questions?



Thanks!

dawnmai@cs.washington.edu

<https://github.com/dawnmai>



www.linkedin.com/in/dawnmai

CREDITS: This presentation template was created by [Slidesgo](#), and includes icons, infographics & images by [Freepik](#)

Please keep this slide for attribution